

# Learnersourcing Quality Assessment of Explanations for *Peer Instruction*

Sameer Bhatnagar<sup>1</sup>, Amal Zouaq<sup>1</sup>, Michel C. Desmarais<sup>1</sup>, and Elizabeth Charles<sup>2</sup>

<sup>1</sup> Ecole Polytechnique Montreal

{sameer.bhatnagar, amal.zouaq, michel.desmarais}@polymtl.ca

<sup>2</sup> Dawson College echarles@dawsoncollege.qc.ca

**Abstract.** This study reports on the application of text mining and machine learning methods in the context of asynchronous peer instruction, with the objective of automatically identifying high quality student explanations. Our study compares the performance of state-of-the-art methods across different reference datasets and validation schemes. We demonstrate that when we extend the task of argument quality assessment along the dimensions of *convincingness*, from curated datasets, to data from a real learning environment, new challenges arise, and simpler vector space models can perform as well as a state-of-the-art neural approach.

**Keywords:** Argument mining · Learnersourcing · Peer Instruction

## 1 Introduction

Learning environments that leverage peer submitted content, carry the great advantage of scaling up to a larger content base. Be it student-generated questions, sample solutions, feedback, or explanations and hints, peer-submitted content can easily help grow a database to meet the demand for practice items and formative assessment. Platforms built on this principle are growing in popularity [7][15], freeing the teacher from the tedious but important task of developing multiple variants of the same items.

Critical to the success of such environments is the capacity to automatically assess the quality of learner-generated content. There are a growing number of tools which put students at the centre of this challenge: students submit their own work, but then are prompted to evaluate a subset of their peers’ submissions, and sometimes even provide feedback [23][3].

To counter the drawback that lies in the varying ability of novices to evaluate and provide good feedback to their peers, these environments often use *pairwise* comparison. It is widely accepted that evaluative judgements, based on ranking two objects relative to one another, are easier to make than providing an absolute score. Adaptive Comparative Judgement [22], where teachers assess student submissions by simply choosing which is better from a pair, has been shown to be a reliable and valid alternative to absolute grading. As such, there is a growing

family of learning tools which, at the evaluation step, present the peer-generated content as *pairs* to the current learner, and prompt for a pairwise ranking. This data can provide feedback to the students who originally submitted the item, but can also be used for moderating learner content.

In platforms where students generate content that is part of the learning activities of future students, filtering out irrelevant and misleading material is paramount. However, while removing bad content is important, educators hope to identify, and subsequently maximize the use of the best student generated items. Since not all students can be asked to evaluate all possible pairs, this in turn leads to the challenge of optimizing which items need evaluation by the “student-come-moderators”, without hindering their learning.

Our work is centred on data coming from the subset of these learning environments that enable *peer instruction*[6], which follow a specific two-stage script:

1. students are prompted to answer a multiple choice question, and provide a free-text explanation that justifies their answer;
2. students are then prompted to reconsider their answer, by presenting them a selection of explanations written by previous students [1]. Students can either decide that their own explanation is best, or indicate which of their peers’ explanation was the most convincing.

We frame the explanations students produce here as *arguments* meant to persuade one’s peers that their own reasoning is best. Recent research in argument mining has proposed a novel machine learning task: given a pair of arguments on a given topic, can we predict which is more *convincing*? [12] [24] [26] [11] Our objective is to evaluate if we can extend this task to a peer-instruction setting: given a pair of explanations, one written by a student, and another by their peer, can we predict which they will select as more convincing?

To our knowledge, this is the first analysis of peer instruction data through the lens of *learnersourcing* and argument *convincingness*. Thus, as a reference, we include datasets and methods from the *argument mining* research community, specifically focused on pairwise preference ranking of arguments for automatic assessment of convincingness. We apply vector space models, as well as a state-of-the-art neural approach, and report on their performance for this task.

Our findings suggest that the arguments generated in learning environments centred on undergraduate science topics present a more challenging variant of the task originally proposed in the argument mining community, and that classical approaches can match neural models for performance on this task across datasets, depending on the context.

## 2 Related Work

### 2.1 Learnersourcing & Comparative Peer Assessment

The term *learnersourcing* has been defined as the process in “which learners collectively generate useful content for future learners, while engaging in a meaningful learning experience themselves” [27]. One of the earliest examples of a learning environment centred on this mode is Peerwise [8], wherein students generate

question items for the subjects they are learning, share them with their peers, so that others can use them for practice. Ripple [15] is a similarly built system, but with an added recommendation engine which adaptively selects which problems to suggest to which students.

Other tools leave the creation of question items to teachers, but call on students to generate and evaluate *explanations* for the answers. The AXIS [28] system prompts students to generate an explanation for their response to a short answer question, and then evaluate a similar explanation from one of their peers on a scale of 1-10 for *helpfulness*. This rating-data drives the reinforcement learning algorithm that decides which explanations to show to future students.

A similar class of learning platforms leverage comparative judgement for supporting how students can evaluate the work of their peers. Juxtappeer [3] asks students to provide feedback to a single peer on their work by explicitly comparing to that of another. ComPAIR [23] asks students for feedback on each item of the pair they are presented with, with a focus on what makes one “better” than the other.

Finally, *peer instruction* question prompts are being used more frequently inside online learning assignments [2] [4]. Moderating the quality of student explanations therein is important work that can begin with unsupervised clustering to filter out irrelevant content[10]; however identifying the *best* explanations that might promote learning, is at the heart of this research.

## 2.2 Argument Quality & Convincingness

Work in the area of automatic evaluation of argument quality finds its roots in detecting evidence in legal texts [18], but has accelerated in recent years as more datasets become available, and focus shifts to modelling more qualitative measures, such as *convincingness*. Some earlier efforts included work on automatic scoring of persuasive essays [21] and modelling persuasiveness in online debate forums [25]. However, evaluating argument *convincingness* with an absolute score can be challenging for annotators, which has led to significant work that starts from data consisting of paired arguments, labelled with which of the two is most convincing.

Research efforts in modelling these pairwise preferences include a feature-rich support vector machine, as well as an end-to-end neural approach based on pre-trained GloVe vectors[20] which are fed into a Bi-directional Long-Short-Term Memory network [12]. Most recently, this work has been extended with a Siamese network architecture, wherein each of the two legs is a BiLSTM network that share weights, trained to detect which argument in a pair has the most convincing evidence[11]. Finally, based on the recent success of transformer models in NLP, the current state of the art for the task of pairwise preference ranking of *convincingness* is based on fine tuning pre-trained weights for the publicly available base BERT [9] model setup for sequence classification[26].

### 3 Data

The objective of this study is to compare and contrast how text mining methods for evaluating argument quality, specifically for argument *convincingness*, perform in an online learning environment with learner-generated and annotated arguments. Our primary dataset comes from a *peer instruction* learning environment myDALITE.org, used by instructors in a context of low-stakes formative assessment in introductory undergraduate science courses. To provide context as to the performance that can be expected for this task in a relatively novel setting, we include in our study three publicly available datasets, each specifically curated for the task of automatic assessment of argument quality along the dimension *convincingness*. Table 1 provides examples of an argument pair from each of the datasets.

#### 3.1 UKP & IBM

**UKPConvArgStrict**[12], hence forth referred to as **UKP**, was the first to propose the task of pairwise preference learning for argument convincingness. The dataset consists of just over 1k individual arguments, that support a particular stance for one of 16 topics, collected from online debate portals and annotated as most/less convincing in a crowd-sourcing platform.

More recently, a second similar dataset was released by the research group associated with IBM Project Debater, **IBMArgQ-9.1kPairs**[26], henceforth referred to as **IBM\_ArgQ**. **IBM\_ArgQ** data is strongly curated with respect to the relative length of the arguments in each pair: in order to control for the possibility that annotators may make their choice of which argument in the pair is more *convincing* based merely on the length of the text, the mean difference in word count,  $\overline{\Delta wc}$ , is just 3 words across the entire dataset, which is 10 times more homogeneous than pairs in **UKP**.

Finally, we include a third reference dataset, **IBM\_Evi** [11]. The important distinction here is that the arguments are actually extracted as evidence for their respective topic from Wikipedia, and hence represent cleaner, more well-formed text than our other reference datasets. We filter this dataset to only include topics that have at least 50 associated argument pairs.

Table 2 summarizes some of the descriptive statistics that can be used to compare these sources, and potentially explain some of our experimental results.

#### 3.2 dalite

The **dalite** dataset is from a *peer instruction* environment explained in the introduction. It contains only the observations where, after writing an explanation for their answer choice, on the review step, students chose a peer’s explanation as more *convincing* than their own.

To ensure internal reliability, we only keep argument explanations that were also chosen by at least 5 different students. To ensure that the explanations in each pair are of comparable length, we keep only those with word counts that are

Table 1: Examples of argument pairs from each dataset. These examples were selected because they were incorrectly classified by all of our models, and demonstrate the challenging nature of the task. In each case, the argument labeled as more convincing is in *italics*.

(a) A pair of arguments from **UKP**, for the prompt topic: “school uniforms are a good idea”.

a1	a2
I take the view that, school uniform is very comfortable. Because there is the gap between the rich and poor, school uniform is efficient in many ways. If they wore to plain clothes every day, they concerned about clothes by brand and quantity of clothes. Every teenager is sensible so the poor students can feel inferior. Although school uniform is very expensive , it is cheap better than plain clothes. Also they feel sense of kinship and sense of belonging. In my case, school uniform is convenient. I don't have to worry about my clothes during my student days.	<i>I think it is bad to wear school uniform because it makes you look unatrel and you cannot express yourself enough so band school uniform OK</i>

(b) A pair of arguments from **IBM\_ArgQ** , for the prompt topic: “We should support information privacy laws”.

a1	a2
<i>if a company is not willing to openly say what they are going to do with my data, they shouldn't be allowed to do it.</i>	if you are against information privacy laws, then you should not object to having a publicly accessible microphone in your home that others can use to listen to your private conversations.

(c) Student explanations from **dalite**, for the question prompt: “Rank the magnitudes of the electric field at point A, B and C shown in the following figure from greatest magnitude to weakest magnitude”.

a1	a2
<i>At B, the electric field vectors cancel (E=0). C is further away than A and is therefore weaker.</i>	A is closest, B experiences the least since it is directly in the middle, and C the least since it is most far away.

Table 2: Descriptive statistics for each dataset of argument pairs, with last rows showing **dalite** data split by discipline.  $N_{\text{args}}$  is the number of individual arguments, distributed across  $N_{\text{pairs}}$  revolving around  $N_{\text{topics}}$ .  $N_{\text{vocab}}$  is the number of unique tokens in all the arguments.  $\overline{wc}$  is the average number of words per argument, shown with the standard deviation ( $SD$ ).  $\overline{\Delta wc}$  is the average relative difference in number of words for each argument in each pair.

dataset	$N_{\text{pairs}}$	$N_{\text{topics}}$	$N_{\text{args}}$	$N_{\text{vocab}}$	$\overline{wc}$ (SD)	$\overline{\Delta wc}$ (SD)
IBM_ArgQ	9125	11	3474	6710	23 ( 7)	3 ( 2)
UKP	11650	16	1052	5170	49 (28)	30 (23)
dalite	8551	102	8942	5571	17 (15)	12 ( 7)
IBM_Evi	5274	41	1513	6755	29 (11)	3 ( 2)
dalite:Biology	3919	49	4116	3170	15 (14)	10 ( 6)
dalite:Chemistry	1666	24	1758	2062	20 (14)	12 ( 7)
dalite:Physics	2966	29	3068	2478	19 (15)	15 ( 7)

within 25 words of each other. This leaves us a dataset with 8551 observations, spanning 2216 learner annotators having completed, on average, 4.0 items each, from a total of 102 items across three disciplines, with at least 50 explanation-pairs per item. We draw an analogy between a “topic” in the reference data, and a “question item” in **dalite**.

Table 3: Number of argument pairs in **dalite**, broken down by discipline, and the correctness of the selected answer choice on initial step, and then on review step.

	rr (right-right)	rw (right-wrong)	wr (wrong-right)	ww (wrong-wrong)
Biology	2459	124	733	603
Chemistry	1151	51	228	236
Physics	2288	66	278	334

Table 3 highlights three key differences between the modelling task of this study, and related work in argument mining.

First, in **IBM\_ArgQ** and **UKP**, annotators are presented pairs of arguments that are always for the same stance, in order to limit bias due to their opinion on the topic when evaluating which argument is more convincing (this is also true of many of the pairs in **IBM\_Evi**). In a *peer instruction* learning environment, other pairings are possible, and pedagogically relevant. In **dalite**, the majority of students keep the same answer choice on the review step, and so they are comparing two explanations that are either both for the correct answer choice (“rr”), or an incorrect answer choice (“ww”). This is analogous to arguments for the *same stance*. However, 17 % of the observations in **dalite** are

for when students not only choose an explanation more convincing than their own, but also switch answer choice, either from the wrong to right, or the reverse (convincingness across *different stances*).

Second, a more fundamental difference is that our study focuses on undergraduate science courses across three disciplines, wherein the notion of answer “correctness” is important. There are a growing number of ethics and humanities instructors using the peer instruction platform, where the question prompts are topics more like a debate, as in our reference datasets. We leave these comparisons for future work.

Thirdly, each argument pair is made up of the one written by the current learner, while the other is an alternative generated by a peer that the current learner chose as more convincing. In this respect, the **dalite** dataset is different than the reference datasets, since it is the same person, the current learner, who is author and annotator. In the reference datasets, data was always independently annotated using crowdsourcing platforms.

## 4 Methodology

Choosing which argument is more convincing from a pair is a binary ordinal regression task, where the objective is to learn a function that, given two feature vectors, can assign the better argument a rank of +1, and the other a rank of -1. It has been proven that such a binary ordinal *regression* problem can be cast into an equivalent binary *classification* problem, wherein the model is trained on the *difference* of the feature vectors of each argument in the pair [13]. Referred to as *SVM-rank*, this method of learning pairwise preferences has been used extensively in the context of information retrieval (e.g. ranking search results for a query) [14], but also more recently in evaluating the journalistic quality of newspaper and magazine articles [17]. The study accompanying the release of **UKP** first proposed this method for modelling argument convincingness [12].

In our study, we explore three models, described below.

**Vector Space Models, ArgBoW:** We follow-up work using SVM-rank, building simple “bag-of-words” vector space models to represent our argument text. We take all of the individual arguments for a particular topic in our training set (known as “explanations” for a particular question item in the case of the **dalite** data), lemmatize the tokens, and build term-document matrices for each topic. We then take the arithmetic difference of these normalized term frequency vector representations of each argument to train Support Vector Machine classifiers on. We refer to this model as **ArgBoW**. We do not, for this study, include any information related to the topic prompt in our document representation.

**Pre-trained word embeddings, ArgGloVe:** A limitation of vector space models for text classification is the exclusion of words that are “out-of-vocabulary” in the test set when compared to the training data. Previous research has addressed this using language models that offer pre-trained word-embeddings that have already learned expansive vocabularies from massive corpora of text [12][11]. In our

**ArgGloVe** model, we encode each token of each argument using 300-dimensional GloVe vectors [20], and represent each argument as the average of its token-level embedding vectors. We then feed these into the same SVM-rank architecture described above.

**Transfer Learning, ArgBERT:** Finally, in order to leverage recent advances in transfer-learning for NLP, the final model we explore is **ArgBERT**. We begin with a pre-trained language model for English built using Bi-directional Encoder Representation from Transformers, known as *BERT*[9], trained on large bodies of text for the task of masked token prediction and sentence-pair inference. As proposed in [26], we take the final 768-dimensional hidden state of the base-uncased BERT model, feed it into a binary classification layer, and fine-tune all of the pre-trained weights for the task of sequence classification using our argument-pair data. As in other applications involving sentence pairs for BERT, each argument pair is encoded as [CLS] A [SEP] B, where the special [SEP] token instructs the model as to the boundary between arguments A and B.<sup>3</sup>

#### 4.1 Baselines & Validation Schemes

We build a baseline model, **ArgLength**, which is trained on simply the number of words in each argument, as there may be many contexts where students will simply choose the longer/shorter argument, based on the prompt.

In order to get a reliable estimate of performance, we employ stratified 5-fold cross-validation for our experiments. This means that for each fold, we train our model on 80% of the available data, ensuring that each topic and output class is represented equally. In the context of peer instruction learning environments, this is meant to give a reasonable estimate of how our models would perform in predicting which explanations will be chosen as most convincing, *after* a certain number of responses have been collected.

However standard practice in the argument mining community is to employ “cross-topic” validation for this task. For each fold, the arguments for one topic (or question item, in the case of **dalite**) are held out from model training. Evaluating model performance on a yet unseen topic is a stronger estimate for how a model will perform when new question items are introduced into the content base. We evaluate our models using both validation schemes.

## 5 Results & Discussion

The performance of each model across the different datasets is presented in figures 1 and 2, using respectively 5-fold cross-validation and cross-topic validation.

For the purposes of comparison, in Table 4, we denote, to the best of our knowledge, the state-of-the-art performance for each dataset on the task of pairwise classification for *convincingness*.

<sup>3</sup> Modified from the `run_glue.py` script provided by the `transformers` package, built by company hugging face. All code for this study provided at <https://github.com/sameerbhatnagar/ectel2020>



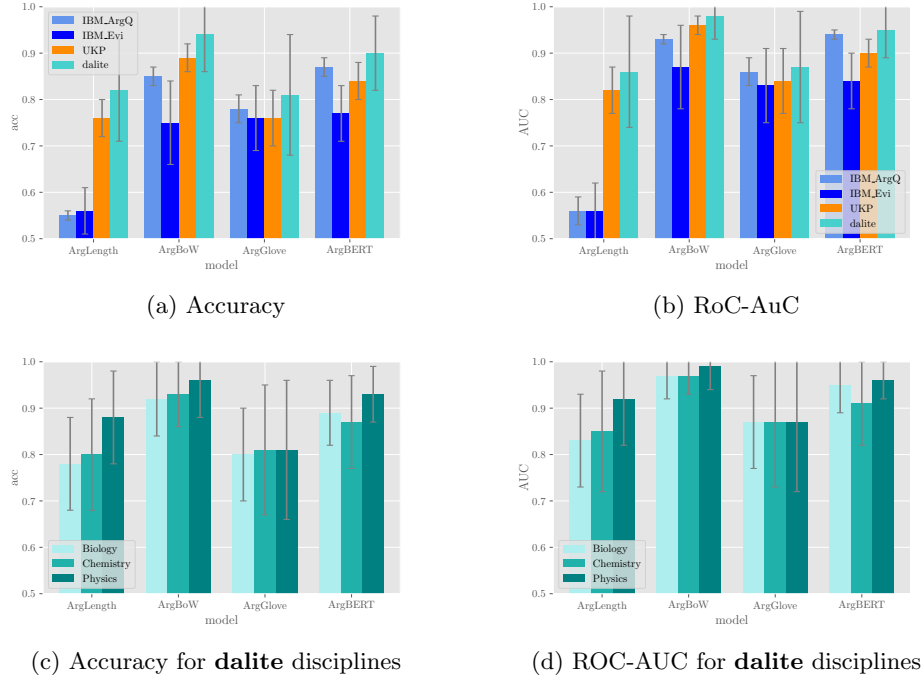


Fig. 1: Pairwise ranking classification accuracy and ROC-AUC, evaluated using **5-fold stratified** cross-validation, for different models across datasets in figures 1a and 1b. Figures 1c and 1d split the performance for the **dalite** dataset across disciplines.

The first point we remark is that the performance of **ArgBERT**, which is the state-of-the-art for our reference datasets, also performs relatively well on **dalite** data. This, in part, provides some support for the premise that student explanations, and the peer-vote-data we extract from a *peer instruction* learning environment, can actually be modelled via a pairwise argument ranking methodology. It also supports our line of inquiry, and lays the foundation for future research in this area.

Second, we see that the baseline performance set by **ArgLength** is not uniform across the datasets, which, we surmise, is due in most part to how carefully curated the datasets are to begin with. **IBM\_ArgQ** and **IBM\_Evi** are datasets that were released, in part, for the purpose of promoting research on automatic assessment of argument quality, and are meant to serve as a benchmark that cannot simply be learned with word counts alone. As can be seen in Table 2,  $\Delta wc$  is greatest in **UKP** and **dalite**, as they are less carefully curated, and contain many more argument pairs that have a larger relative difference in length. This is particularly relevant in the context of learning environments based on *peer instruction*: depending on the importance placed on the review step by their

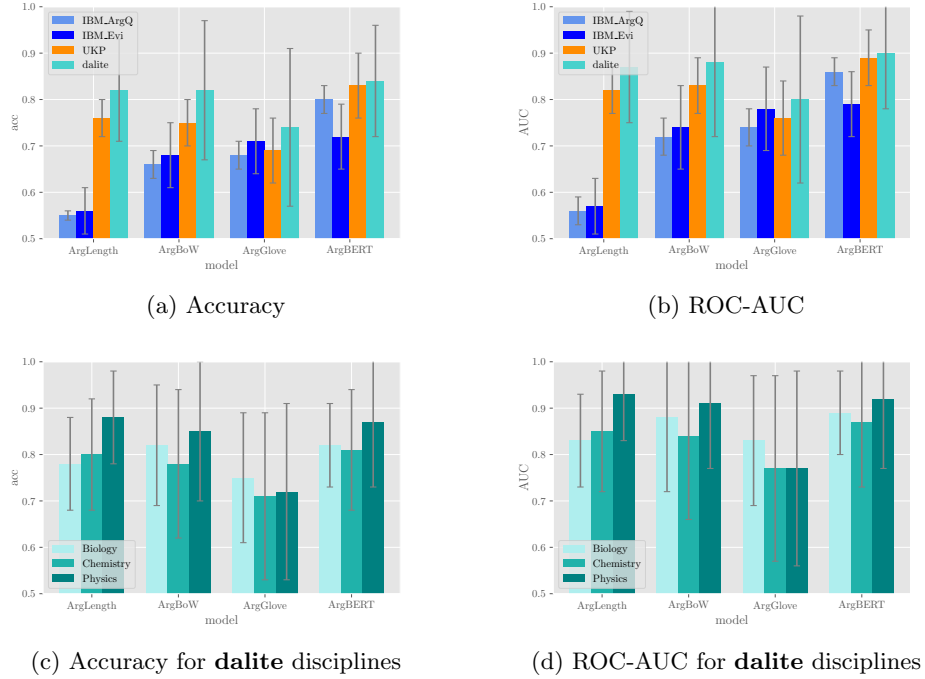


Fig. 2: Pairwise ranking classification accuracy and ROC-AUC, evaluated using **cross-topic** validation, for different models across datasets in figures 2a and 2b. Figures 2c and 2d split the performance for the **dalite** dataset across disciplines.

teacher, students may, or may not, truly engage with their peers’ explanations, reverting to simply choosing explanations based solely on how many words they have to read (or not). This sets the bar very high for other more sophisticated approaches.

Third, we present the results for two different validation schemes, in Figures 1 and 2, to highlight the impact they may have on model selection. When evaluating performance using 5-fold cross-validation (figure 1), where some topic data is held-in for the training, the relatively simple **ArgBoW** performs at least as well, if not better, than all other models, including the state-of-the-art **ArgBERT**. This is not the case in figure 2, where each individual topic is completely held out from training (standard practice in argument mining research). We surmise that **ArgBoW** models can learn the necessary requisite vocabulary to explain what makes an argument more *convincing* and perform well under conditions where at least some data is available for a topic. This effect is pronounced for **dalite**, where the vocabulary is relatively constrained, as seen in Table 2, where the ratio of  $N_{\text{vocab}}$  to  $N_{\text{args}}$  is lowest for **dalite**.

This result is not without precedent: in a study on the task of pairwise ranking of newspaper articles based on “quality”, the authors achieve a similar

Table 4: State of the art performance for pairwise argument classification of convincingness for three publicly available datasets, using cross-topic validation scheme

Dataset	Acc	AUC	model
UKP	0.83	0.89	ArgBERT[26]
IBM_ArgQ	0.80	0.86	ArgBERT[26]
IBM_Evi	0.73	-	EviConvNet[11]

result: when comparing the performance of SVM-rank models using different input feature sets (e.g. *use of visual language*, *use of named entities*, *affective content*), their top performing models achieve pairwise ranking accuracy of 0.84 using a combination of content and writing features, but also a 0.82 accuracy with the content words as features alone[17]. While **ArgBoW** will suffer from the *cold-start* problem when new question items are added to the set of learning activities, as no student explanations are yet available, and the vocabulary is still unknown, this may be remedied by the addition of calculated features to the model.

However the more fundamental result here may point to the importance of verifying multiple validation schemes in future work that ties together methods and data from learning analytics and argument mining research. Cross-topic validation indicates that **ArgBERT** performs better than **ArgBoW**, likely because the model does not overfit to the training data, and the final hidden state of the transformer has learned something more general about the task. Yet even better results might be obtained with a simpler vector space model after some student answers have been collected.

Fourth, we observe the performance of models across different disciplines in **dalite**. Results seem to indicate that argument pairs from items in **dalite:Physics** are easiest to classify. This effect is even more important when we take into account that we have the least data from this discipline ( $N_{\text{pairs}}$  in Table 2). This may be in part due to the impact of upstream tokenizers, which fail to adequately parse arguments that have a lot of chemical equations, or molecular formulae. Most of the items in **dalite:Physics** are conceptual in nature, and the arguments contain fewer non-word tokens.

Finally, we highlight the variance in performance across models, datasets, and validation schemes. We posit that the task of pairwise classification of argument quality, along the dimension of *convincingness*, is more challenging when the data is generated by students as part of a learning activity, than with data collected from crowd-sourcing annotation platforms and online debate portals. This effect is expected to be more pronounced in explanations for items from STEM education, as the difference between correct and incorrect answer choices will be more prevalent than in learning contexts focused on humanistic disciplines. When a student is comparing their explanation, which may be for a incorrect answer choice, with that for an explanation of a correct answer choice, we refer to this as “wr”. This would be the equivalent of argument pairs which contained

arguments of opposite stance, which is only true in **IBM\_Evi**. Of note is the relative stable performance of **ArgGlove** across datasets. This may further indicate that there is promise in vector space approaches, as the semantic information captured in well-trained word embeddings can be leveraged to address the challenge when there is large variance in words used to express arguments (most pronounced in **UKP**).

## 6 Conclusion and Future Work

Asynchronous peer instruction offers the potential of not only *learnersourcing* explanations for formative assessment items, but the curation of those explanations as well. This is a scalable modality that enables students to interact with their peers' contributions, all as part of the learning process. However, to our knowledge, we are among the first to explore methods that aim to distill the data from such a context and automatically assess the quality of student explanations along the dimension of *convincingness*. The implications of this study for technology-enhanced-learning are in the design of peer instruction platforms. A machine learning approach, with careful choice of validation scheme, can help automatically identify the explanations students find most convincing. Directions for future work lie in exploring the features that these models leverage for their predictions, and developing frameworks for providing feedback to both teachers and students.

The performance of vectors space models in this study indicate that more work should be done in expanding the representation of text arguments using calculated features. This approach has been the best performing in other studies [17,19], where a combination of writing and content features to achieve the best results. This avenue must be explored more thoroughly, especially as this may vary across disciplines and teaching contexts.

This current study only addressed the task of predicting which of the explanations was more convincing in pair. The next task for future work lies in *Learning to Rank*, wherein this pairwise preference data is used to infer a global ranking of argument quality.

## 7 Acknowledgements

Funding for the development of myDALITE.org is made possible by *Entente-Canada-Quebec*, and the *Ministère de l'Éducation et Enseignement Supérieure du Québec*. Funding for this research was made possible by the support of the Canadian Social Sciences and Humanities Research Council *Insight* Grant. This project would not have been possible without the SALTISE/S4 network of researcher practitioners, and the students using myDALITE.org who consented to share their learning traces with the research community.

## References

1. Bhatnagar, S., Lasry, N., Desmarais, M., Charles, E.: DALITE: Asynchronous Peer Instruction for MOOCs. In: European Conference on Technology Enhanced Learning. pp. 505–508. Springer (2016)
2. University of British Columbia, T.L.T.: ubc/ubcpi (Aug 2019), <https://github.com/ubc/ubcpi>, original-date: 2015-02-17T21:37:02Z
3. Cambre, J., Klemmer, S., Kulkarni, C.: Juxtapeer: Comparative Peer Review Yields Higher Quality Feedback and Promotes Deeper Reflection. In: Proceedings of the 2018 CHI Conference on Human Factors in Computing Systems - CHI '18. pp. 1–13. ACM Press, Montreal QC, Canada (2018). <https://doi.org/10.1145/3173574.3173868>
4. Charles, E.S., Lasry, N., Bhatnagar, S., Adams, R., Lenton, K., Brouillette, Y., Dugdale, M., Whittaker, C., Jackson, P.: Harnessing peer instruction in- and out-of class with myDALITE. In: Fifteenth Conference on Education and Training in Optics and Photonics: ETOP 2019. p. 11143–89. Optical Society of America (2019), [http://www.osapublishing.org/abstract.cfm?URI=ETOP-2019-11143\\_89](http://www.osapublishing.org/abstract.cfm?URI=ETOP-2019-11143_89)
5. Chen, X., Bennett, P.N., Collins-Thompson, K., Horvitz, E.: Pairwise ranking aggregation in a crowdsourced setting. In: Proceedings of the sixth ACM international conference on Web search and data mining. pp. 193–202 (2013)
6. Crouch, C.H., Mazur, E.: Peer instruction: Ten years of experience and results. *American Journal of Physics* **69**(9), 970–977 (2001)
7. Denny, P.: The Effect of Virtual Achievements on Student Engagement. In: Proceedings of the SIGCHI Conference on Human Factors in Computing Systems. pp. 763–772. CHI '13, ACM, New York, NY, USA (2013). <https://doi.org/10.1145/2470654.2470763>, <http://doi.acm.org/10.1145/2470654.2470763>, event-place: Paris, France
8. Denny, P., Hamer, J., Luxton-Reilly, A., Purchase, H.: PeerWise: Students Sharing Their Multiple Choice Questions. In: Proceedings of the Fourth International Workshop on Computing Education Research. pp. 51–58. ICER '08, ACM, New York, NY, USA (2008). <https://doi.org/10.1145/1404520.1404526>, <http://doi.acm.org/10.1145/1404520.1404526>, event-place: Sydney, Australia
9. Devlin, J., Chang, M.W., Lee, K., Toutanova, K.: Bert: Pre-training of deep bidirectional transformers for language understanding. arXiv preprint arXiv:1810.04805 (2018)
10. Gagnon, V., Labrie, A., Desmarais, M., Bhatnagar, S.: Filtering non-relevant short answers in peer learning applications. In: Proc. Conference on Educational Data Mining (EDM) (2019)
11. Gleize, M., Shnarch, E., Choshen, L., Dankin, L., Moshkovich, G., Aharonov, R., Slonim, N.: Are you convinced? choosing the more convincing evidence with a Siamese network. arXiv preprint arXiv:1907.08971 (2019), <https://www.aclweb.org/anthology/P19-1093/>
12. Habernal, I., Gurevych, I.: Which argument is more convincing? Analyzing and predicting convincingness of Web arguments using bidirectional LSTM. In: Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers). pp. 1589–1599. Association for Computational Linguistics, Berlin, Germany (Aug 2016). <https://doi.org/10.18653/v1/P16-1150>, <https://www.aclweb.org/anthology/P16-1150>
13. Herbrich, R., Graepel, T., Obermayer, K.: Support vector learning for ordinal regression (1999), publisher: IET

14. Joachims, T.: Optimizing search engines using clickthrough data. In: Proceedings of the eighth ACM SIGKDD international conference on Knowledge discovery and data mining. pp. 133–142 (2002)
15. Khosravi, H., Kitto, K., Williams, J.J.: Ripple: A crowdsourced adaptive platform for recommendation of learning activities. arXiv preprint arXiv:1910.05522 (2019)
16. Lippi, M., Torroni, P.: Argumentation mining: State of the art and emerging trends. *ACM Transactions on Internet Technology (TOIT)* **16**(2), 10 (2016)
17. Louis, A., Nenkova, A.: What Makes Writing Great? First Experiments on Article Quality Prediction in the Science Journalism Domain. *Transactions of the Association for Computational Linguistics* **1**, 341–352 (2013)
18. Moens, M.F., Boiy, E., Palau, R.M., Reed, C.: Automatic detection of arguments in legal texts. In: Proceedings of the 11th international conference on Artificial intelligence and law. pp. 225–230 (2007)
19. Nguyen, D., Doğruöz, A.S., Rosé, C.P., de Jong, F.: Computational Sociolinguistics: A Survey. arXiv preprint arXiv:1508.07544 (2015)
20. Pennington, J., Socher, R., Manning, C.D.: Glove: Global Vectors for Word Representation. In: EMNLP. vol. 14, pp. 1532–1543 (2014)
21. Persing, I., Ng, V.: End-to-End Argumentation Mining in Student Essays. In: HLT-NAACL. pp. 1384–1394 (2016)
22. Pollitt, A.: The method of Adaptive Comparative Judgement. *Assessment in Education: Principles, Policy & Practice* **19**(3), 281–300 (2012). <https://doi.org/10.1080/0969594X.2012.665354>, <https://doi.org/10.1080/0969594X.2012.665354>, publisher: Routledge eprint: <https://doi.org/10.1080/0969594X.2012.665354>
23. Potter, T., Englund, L., Charbonneau, J., MacLean, M.T., Newell, J., Roll, I., others: ComPAIR: A new online tool using adaptive comparative judgement to support learning with peer feedback. *Teaching & Learning Inquiry* **5**(2), 89–113 (2017)
24. Simpson, E., Gurevych, I.: Finding Convincing Arguments Using Scalable Bayesian Preference Learning. *Transactions of the Association for Computational Linguistics* **6**, 357–371 (2018), <https://www.aclweb.org/anthology/Q18-1026>
25. Tan, C., Niculae, V., Danescu-Niculescu-Mizil, C., Lee, L.: Winning Arguments: Interaction Dynamics and Persuasion Strategies in Good-faith Online Discussions. arXiv:1602.01103 [physics] pp. 613–624 (2016). <https://doi.org/10.1145/2872427.2883081>, <http://arxiv.org/abs/1602.01103>, arXiv: 1602.01103
26. Toledo, A., Gretz, S., Cohen-Karlik, E., Friedman, R., Venezian, E., Lahav, D., Jacovi, M., Aharonov, R., Slonim, N.: Automatic Argument Quality Assessment-New Datasets and Methods. In: Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP). pp. 5629–5639 (2019), <https://www.aclweb.org/anthology/D19-1564.pdf>
27. Weir, S., Kim, J., Gajos, K.Z., Miller, R.C.: Learnersourcing subgoal labels for how-to videos. In: Proceedings of the 18th ACM Conference on Computer Supported Cooperative Work & Social Computing. pp. 405–416. ACM (2015)
28. Williams, J.J., Kim, J., Rafferty, A., Maldonado, S., Gajos, K.Z., Lasecki, W.S., Heffernan, N.: AXIS: Generating Explanations at Scale with Learnersourcing and Machine Learning. In: Proceedings of the Third (2016) ACM Conference on Learning @ Scale - L@S '16. pp. 379–388. ACM Press, Edinburgh, Scotland, UK (2016). <https://doi.org/10.1145/2876034.2876042>